

State-based Endorsement

(formerly known as “ownable state”, also known as “key-level validation”)


M. Neugschwandtner, A. Sorniotti

Motivation

- Endorsement policies are currently
 - coarse-grained, on a per-chaincode level
 - change requires chaincode upgrade
- Proposal: state-based endorsement
 - endorsement policies on a per-KVS-key-level
 - changes require only chaincode logic

State-based endorsement

- KVS key with designated endorsement policy
 - EP encoded as key-level metadata
 - EP used to validate changes to KVS value



Comparison

	Chaincode Endorsement Policy	State-based Endorsement Policy
Modification process	tied to the the chaincode lifecycle: chaincode upgrade required to set a new endorsement policy	part of a regular fabric transaction that updates key metadata
Granularity	one endorsement policy per chaincode: different endorsement policies require multiple chaincodes	one endorsement policy per KVS key

Chaincode Shim Interface


```
// SetStateEndorsementPolicy sets the endorsement policy for the specified
// `key`. The endorsement policy can not be set on a key that does not
// exist. To remove the endorsement policy for the given key, set `ep` to nil.
SetStateEndorsementPolicy(key string, ep []byte) error

// GetStateEndorsementPolicy retrieves the endorsement policy for `key`.
// Note that this will introduce a read dependency on `key`.
GetStateEndorsementPolicy(key string) ([]byte, error)

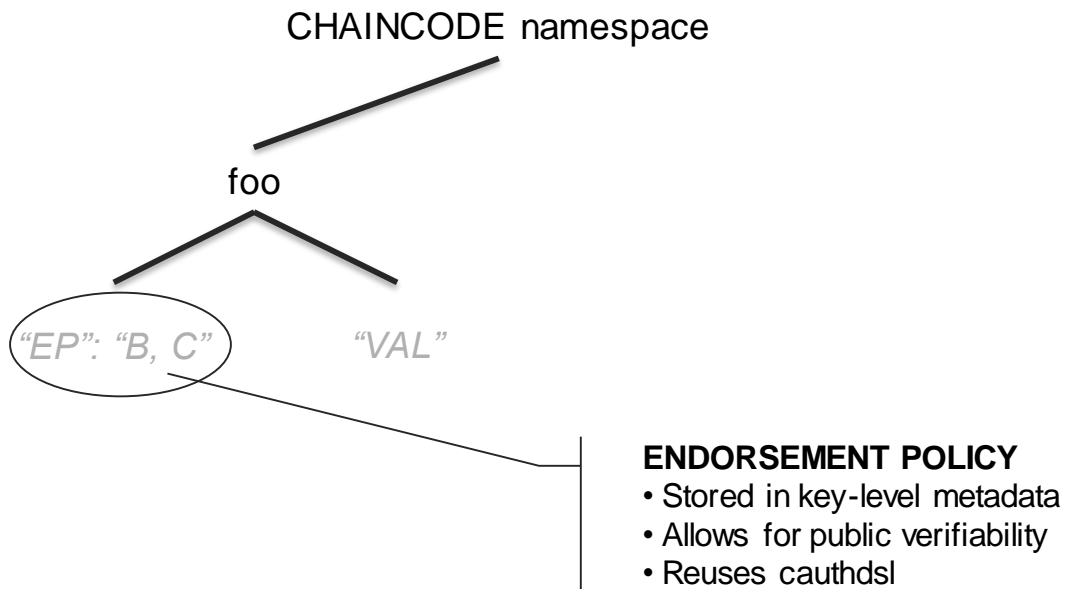
// SetPrivateDataEndorsementPolicy sets the endorsement policy for the private
// data specified by `key`.
SetPrivateDataEndorsementPolicy(collection, key string, ep []byte) error

// GetPrivateDataEndorsementPolicy retrieves the endorsement policy for the
// private data specified by `key`. Not that this introduces a read dependency
// on `key`.
GetPrivateDataEndorsementPolicy(collection, key string) ([]byte, error)
```


On Ledger Representation




On Ledger Representation



On Ledger Representation



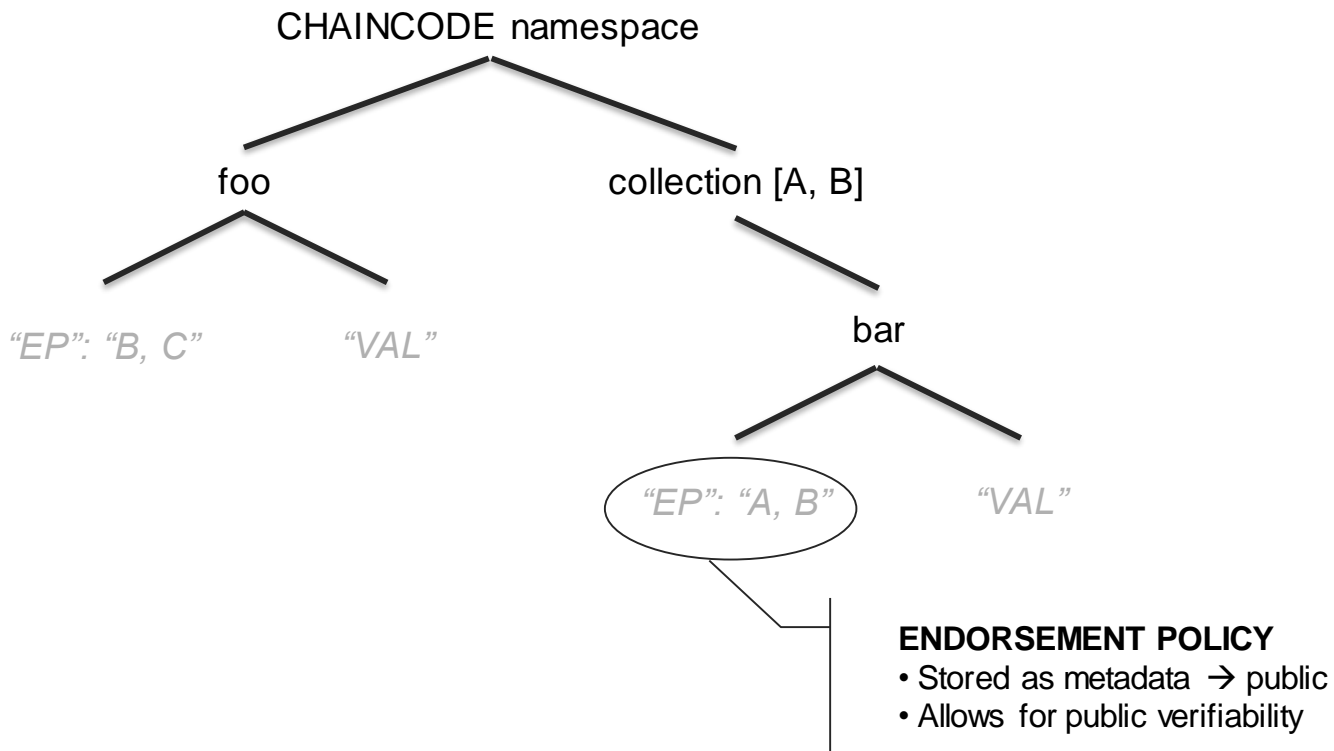
On Ledger Representation




COLLECTION

- Can be used together with key-level metadata
- Way to control both dissemination and endorsement
- Dissemination policy must be superset of EP


On Ledger Representation



Workflow




Workflow




`S.METADATA[EP]` → red, green

Workflow





`S.METADATA[EP]` →

Workflow




S.METADATA[EP] →


Workflow



Workflow



1. Lookup $S.METADATA[EP]$
2. Verify that tx complies with policy therein
3. If so, commit; else, abort



Integration with plugins

- State-based endorsement will be delivered as a default fabric feature
- Default validation plugin will be extended to support key-level endorsements
- No change required to endorsement plugin
- Validation plugin must perform key-level endorsement policy lookup for every key
 - Even for keys that never had an endorsement policy set
 - (Possible) ledger optimization: avoid trip to DB and return nil endorsement policy for keys that never had a key-level EP set

Validation Rules

Validation	Key-level validation parameter on the ledger	Key-level validation parameter in the rwset
Standard case: check against chaincode EP	no	no
Creation of key-level EP: check against chaincode EP	no	yes
Value update for key with key-level EP: check against key-level EP in the ledger	yes	no
State-level EP update: check against key-level EP in the ledger	yes	yes

NOTE: it is now possible that a tx is committed without a check against the cc EP!

Leder metadata interface

```
type QueryExecutor interface {
    ...
    // GetStateMetadata returns the metadata for given namespace and key
    GetStateMetadata(namespace, key string) (map[string][]byte, error)
    // GetPrivateDataMetadata gets the metadata of a private data item
    GetPrivateDataMetadata(namespace, collection, key string) (map[string][]byte, error)
    ...
}

type TxSimulator interface {
    QueryExecutor
    ...
    // SetStateMetadata sets the metadata associated with an existing key-tuple
    SetStateMetadata(namespace, key string, metadata map[string][]byte) error
    // DeleteStateMetadata deletes the metadata (if any) associated with an existing key-tuple
    DeleteStateMetadata(namespace, key string) error
    // SetPrivateDataMetadata sets the metadata associated with an existing key-tuple
    SetPrivateDataMetadata(namespace, collection, key string, metadata map[string][]byte) error
    // DeletePrivateDataMetadata deletes the metadata associated with an existing key-tuple
    DeletePrivateDataMetadata(namespace, collection, key string) error
    ...
}
```


Metadata behaviour

- Endorsement policy is stored as ledger key metadata
 - Endorsement policy is currently only user of this facility
- Metadata is a map[mkey, mvalue] per KVS key
 - Current interface only permits full map (over)write
- Metadata can't be set for non-existing KVS pair
- Tx can
 - Modify data only
 - Current metadata value will be kept
 - Modify metadata only
 - Current data value will be kept
 - Modify data and metadata
 - Delete metadata
 - Current data value will be kept
 - Delete KVS key
 - Both data and metadata deleted
- Only one read dependency exists, irrespective of whether only data, only metadata or both were read

Capabilities and fork-resistance

- In a network with some 1.3 peers and some <1.3 peers, a chaincode is executed on a 1.3 peer and it performs a `SetStateEndorsementPolicy` call for key “foo”. Tx is ordered and committed
 - 1.3 peer
 - Sets EP on the ledger
 - Checks future changes of “foo” against key-level EP
 - <1.3 peer
 - Ignores metadata `rwset`
 - Doesn't perform any special validation
- New capability introduced and used
 - In the peer-side of the shim to ensure metadata can be set/retrieved only in a 1.3 network
 - In the standard validation plugin to perform key-level validation only in a 1.3 network

Special Case in Validation



- When Tx j is validated by the validation plugin, the status of Tx i is undetermined
 - Due to parallelism, validation of Tx i might be scheduled after that of Tx j
 - MVCC validation takes place sequentially afterwards

To ensure correctness..

- if Tx i is valid according to the validation plugin, invalidate Tx j with a new error code (like we do today for cc upgrade)
- Create lock structure to ensure that
 - Parallel validation occurs to the extent possible
 - Create lock structure to ensure validation of Tx j only starts when validity status of Tx i is known

Thank you